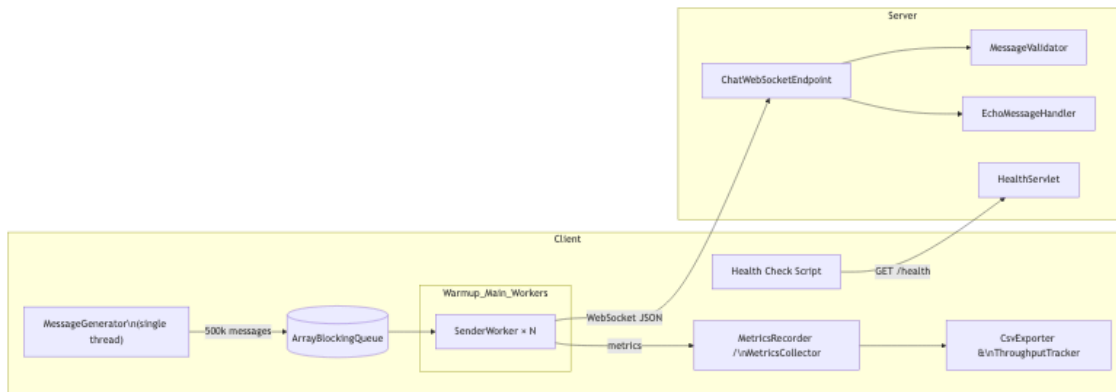


CS6650 Assignment 1 – Design Document

Git repository: https://github.com/VivianDongChen/cs6650_assignment.git

Modules: `/server` , `/client-part1` , `/client-part2` , `/results`

1. Architecture Overview



The client simulates load in two phases (warmup + main) using a shared message queue. Each sender maintains its own WebSocket connection to the Tomcat-hosted server. Metrics are captured locally and exported for analysis (Part 2/Part 3).

EC2 Deployment Evidence:

The screenshot displays the AWS Management Console for an EC2 instance. The instance is named 'i-08e7be9f79d132a46' and is in the 'Running' state. It is a 't3.micro' instance type. The console shows the following details:

- Instance summary:** Instance ID, Name, IP address, Hostname type, Answer private resource DNS name, Auto-assigned IP address, IAM Role, IMDSv2, Operator.
- Public IPv4 address:** 44.250.41.12
- Private IPv4 address:** 172.31.34.223
- Public DNS:** ec2-44-250-41-12.us-west-2.compute.amazonaws.com
- Elastic IP addresses:** None
- AWS Compute Optimizer finding:** Opt-in to AWS Compute Optimizer for recommendations.
- Auto Scaling Group name:** None
- Managed:** false
- Platform details:** Linux/UNIX

2. Major Classes and Relationships

Server Module (/server)

Class	Responsibility	Key Collaborators
ChatWebSocketEndpoint	Entry point for /chat/{roomId} WebSocket sessions. Parses payloads, delegates validation, and echoes results.	MessageValidator, EchoMessageHandler
MessageValidator	Validates userId, username, message, timestamp, messageType against assignment rules.	ChatMessage
EchoMessageHandler	Builds success/error JSON responses with server timestamps.	ChatMessage
ChatMessage	Jackson-mapped POJO for inbound payloads.	Used by endpoint + validator
HealthServlet	/health REST endpoint returning {status,timestamp} JSON for monitoring.	ObjectMapperProvider
ObjectMapperProvider	Centralizes Jackson configuration (JavaTimeModule).	All JSON consumers

Client Part 1 (/client-part1)

Class	Responsibility	Key Collaborators
App	Bootstraps client, loads config, wires generator + orchestrator.	ClientConfig, MessageGenerator, SenderOrchestrator
ClientConfig / ClientConfigLoader	Immutable configuration + CLI parsing with defaults (500k messages, warmup 32x1000, etc.).	All client components
MessageGenerator	Single thread generating 500k messages with required distributions into a bounded queue.	ChatMessage, queue
SenderOrchestrator	Manages warmup/main phases, submits SenderWorker tasks to executor, records timing.	SenderWorker, MetricsRecorder
SenderWorker	Maintains a persistent WebSocket connection, sends messages with retries (max 5) and exponential backoff.	ReliableWebSocketClient, MetricsRecorder
ReliableWebSocketClient	Thin wrapper over java.net.http.WebSocket providing connect/reconnect/send-and-wait semantics.	SenderWorker

MetricsRecorder	Thread-safe counters for successes, failures, retries, connections, runtime, throughput.	SenderOrchestrator, SenderWorker
-----------------	--	-------------------------------------

Client Part 2 (/client-part2)

Extends Part 1 with additional metrics & exports:

Class	Responsibility	Notes
MessageMetric, MetricsCollector	Per-message latency capture and thread-safe aggregation.	Stores 500k records in memory
StatisticsAnalyzer, Statistics	Compute mean/median/p95/p99/min/max, message type & room distribution.	Used after run completion
ThroughputTracker	Buckets acknowledgements into 10s windows for throughput CSV/chart.	
CsvExporter	Writes metrics and throughput CSV outputs to results/part2-output.	
MetricsRecorder (extended)	Includes per-phase timing hooks and integrates with collectors.	

All Part 2 classes reuse Part 1 networking/threading components to keep behavior consistent.

3. Threading Model

Message Generation

- **Single producer thread** (MessageGenerator) pre-generates up to 500,000 messages.
- Uses ArrayBlockingQueue (capacity 10,000) to buffer messages; workers block if queue drains.
- Messages follow spec distributions: `userId` 1–100,000, `username` = `user<id>` , 50 message templates, `roomId` 1–20, `messageType` mix 90% TEXT / 5% JOIN / 5% LEAVE, `timestamp` = `Instant.now()` .

Warmup Phase

- Exactly 32 threads, each sending 1,000 messages (32,000 warmup messages).
- Workers establish new WebSocket connections; metrics captured separately.

Main Phase

- Configurable thread count (typically 128 locally, 224–320 in EC2 tests).
- SenderOrchestrator divides remaining messages evenly across workers.
- Executors use cached thread pool; each worker runs to completion (no reuse).

Metrics Threads (Part 2)

- Metrics collection is non-blocking: workers record metrics and return to sending.
- Post-run aggregation (single thread) generates statistics and CSV outputs.

4. WebSocket Connection Management

1. Connection Lifecycle

- Each `SenderWorker` owns a `ReliableWebSocketClient`.
- On start, `connect()` is invoked (timeout configurable, default `PT5S`).
- After sending all assigned messages, the worker closes the connection gracefully.

2. Retry Semantics

- For each message: up to 5 attempts (`max-retries=5`).
- Backoff starts at `initial-backoff` (default 50ms) and doubles until `max-backoff` (default 200ms Part 1, 1s Part 2).
- On retry, `metrics.addRetries(1)` is recorded.

3. Reconnection Strategy

- Any failure during send triggers `client.reconnect()` (closing + re-opening).
- Successful reconnection increments `metrics.incrementConnectionReconnected()`.
- Certain failures (e.g., invalid payload) mark message as failed after final attempt.

4. Thread Safety

- `ReliableWebSocketClient` synchronizes critical sections to avoid concurrent writes on a single WebSocket.
- Message queue + metrics classes use `BlockingQueue` and `AtomicLong` / `ConcurrentLinkedQueue` to avoid contention.

5. Server Handling

- `ChatWebSocketEndpoint` validates messages and echoes them with server timestamps.
- Invalid payloads return `{status:"error",errors:[...]}`; valid payloads include `serverReceiveTime` and `serverSendTime`.
- Endpoint logs errors and performs best-effort close on IO exceptions.

5. Little's Law Calculations & Observations

Little's Law: $\lambda = L / W$ (throughput equals concurrency divided by service time).

Local Environment (Loopback)

- **Single-thread baseline:** 1,000 messages in **1,916 ms** → `W_local = 1.916 ms`.
- Predicted vs observed throughput:

Threads (L)	Predicted $\lambda = L / W_{\text{local}}$ (msg/s)	Observed λ (msg/s)	Efficiency
48	$\approx 25,052$	16,061	64.1%
64	$\approx 33,403$	20,435	61.2%
96	$\approx 50,104$	27,382	54.7%
128	$\approx 66,805$	30,255	45.3%
160	$\approx 83,507$	28,249	33.8%

Takeaway: Optimal throughput occurs at 128 threads (≈30k msg/s). Efficiency drops as contention and retries grow, revealing diminishing returns beyond 128 threads.

EC2 Environment (Public Internet)

- **Single-thread baseline:** 1,000 messages in **17,371 ms** → $W_{ec2} = 17.371\text{ ms}$ (≈15.455 ms slower than local).
- Predicted vs observed throughput:

Threads (L)	Predicted $\lambda = L / W_{ec2}$ (msg/s)	Observed λ (msg/s)	Efficiency
224	≈ 12,900	6,747	52.3%
256	≈ 14,739	7,687	52.2%
288	≈ 16,582	7,941	47.9%
320	≈ 18,424	5,325	28.9%

Takeaway: Higher latency necessitates more concurrency (288 threads) to hide network delays. Efficiency remains ~50% until overload at 320 threads causes timeouts and failures.

Part 2 Enhancements

- Detailed metrics (latency percentiles, throughput buckets) confirm Little’s Law insights:
 - Mean latency ≈ 21.4 ms, aligning with predicted 17.371 ms + retry overhead.
 - Peak throughput 10,972 msg/s in Part 2 indicates improved connection reuse and instrumentation overhead is negligible.

6. Deployment & Testing Notes

- **Server Deployment:** Tomcat 9 on AWS EC2 (us-west-2). Health verified via `/health` and WebSocket echo tests (`wscat` , Postman).
- **Client Execution:** Maven builds shaded JARs for Part 1 & Part 2. CLI flags allow tuning thread counts, retries, and URIs.
- **Results Artifacts:** Screenshots, CSV exports, charts, and this design document reside in `/results/` for quick reference.
- **EC2 Evidence:** The included AWS console screenshot (see above) documents the running instance (region us-west-2, public IP 44.250.41.12) used for performance testing.